

Project Description: IoT Data Streaming, Local & Global Model Training, and Analytics Pipeline

Project Overview

This project aims to build a complete IoT data processing and analytics pipeline that simulates **live streaming from IoT devices**, performs **per-device local model training**, aggregates these into a **global model**, and evaluates data with **batch and streaming analytics**. The system is designed to be **modular, scalable, and real-time**, allowing seamless addition of devices and expansion of analytics capabilities.

Project Objectives

1. **Simulate live IoT data ingestion** using an open-source dataset.
 2. **Train local models per IoT device** to capture device-specific patterns.
 3. **Aggregate local models into a global model** using federated learning principles.
 4. **Evaluate batch and streaming data** using Apache Spark with the global model.
 5. **Store and manage all data, models, and results** efficiently in MongoDB.
 6. **Visualize per-device and global analytics** using Superset.
-

Data Source

- The project will use **Intel Lab Data**, a public IoT sensor dataset containing readings from multiple devices over time (temperature, humidity, light, voltage).

- The dataset is downloaded and preprocessed to simulate real-time streaming of sensor readings.
-

System Architecture

The system is composed of the following components:

1. Kafka – Live Data Streaming

- **Purpose:** Serve as the backbone for real-time data flow.
- **Functionality:**
 - Read rows from the dataset and send each row as a Kafka message to the topic `iot_stream`.
 - Each message contains:
 - `device_id`
 - Timestamp
 - Sensor readings (temperature, humidity, light, voltage)
 - **Loop over the dataset repeatedly:** once the dataset ends, start again from the beginning to maintain a continuous stream.
- **Output:** Kafka topic `iot_stream` for consumption by Flink and Spark.

2. Apache Flink – Per-Device Processing and Local Modeling

- **Purpose:** Process live streaming data and maintain device-specific models.
- **Functionality:**
 - Consume messages from Kafka topic `iot_stream`.
 - Group data by `device_id`.

- For each device:
 - Maintain a **local neural network**.
 - Perform **feedforward, backpropagation, and weight updates** continuously.
 - Send periodic model updates to the global model server.
- Dynamically handle **new devices**, creating models as needed.
- **Output:** Local model updates sent to the global server.

3. Local Neural Network Models

- **Purpose:** Capture patterns and anomalies specific to each device.
- **Structure:** Lightweight feedforward neural networks.
- **Training:** Incremental training on live streaming data.
- **Storage:** Local models are stored in `/models/local` and updates are sent to the global model server.

4. Global Model Server

- **Purpose:** Aggregate local models into a single global model.
- **Functionality:**
 - Collect model updates from local nodes periodically.
 - Perform federated aggregation to update the global neural network weights every half hour.
 - Maintain a **historical record of updates** with timestamps.
 - Provide the global model to Spark for batch and streaming evaluations.
- **Storage:** Global model stored in `/models/global` and MongoDB.

5. Apache Spark – Batch and Streaming Analytics

- **Batch Analysis:**
 - Read the full dataset to perform heavy analytics.
 - Evaluate predictions using the **latest global model**.
- **Streaming Analysis:**
 - Consume Kafka topic `iot_stream`.
 - Evaluate predictions in near real-time using the global model.
- **Output:** Store both batch and streaming predictions in MongoDB for visualization.

6. MongoDB – Data Storage

- **Collections:**
 - `local_models`: Updates from per-device models.
 - `global_model`: Aggregated global model weights.
 - `predictions`: Batch and streaming evaluation results.
- **Purpose:** Centralized storage to support analytics and visualization.

7. Visualization

- **Tool:** Superset
- **Visualizations:**
 - Per-device sensor readings and predictions.
 - Global model evaluation metrics.
 - Comparison between streaming and batch results.
 - Any other helpful insights

Data Flow

1. Dataset is **downloaded and preprocessed**.
 2. Kafka streams each row of the dataset as a message to the topic `iot_stream`, and automatically loops back to the beginning once the dataset ends — continuously simulating a real IoT data stream since no live data source is available.
 3. Flink **consumes messages**, trains local models, and sends updates to the global model server.
 4. Global server **aggregates local updates** periodically, updates the global model, and stores it in MongoDB.
 5. Spark **performs batch and streaming evaluation** using the global model and stores predictions in MongoDB.
 6. Superset connects to MongoDB to **visualize results** for monitoring and analysis.
-

Project Structure

```
IoT-Streaming-ML-Pipeline/
|
├─ data/
|   └─ raw/
|       └─ intel_lab_data.csv           # Original downloaded
dataset
|   └─ processed/
|       └─ processed_iot_data.csv       # Cleaned &
preprocessed version
|   └─ download_dataset.py              # Script to
automatically download the dataset
|       └─ preprocess_dataset.py        # Script to clean &
prepare data
```

```

|
|— kafka/
|   |— kafka_producer.py           # Reads dataset rows &
sends to Kafka topic
|   |— kafka_consumer_test.py      # Simple test consumer
(for debugging)
|   |— docker-compose.yml          # Kafka + Zookeeper
setup
|   |— kafka_setup.sh              # Helper script to
create topics & start broker
|       |— config/
|           |— kafka_config.json   # Topic names, broker
addresses, configs
|
|— flink/
|   |— flink_job.py                # Main Flink streaming
job (consumes Kafka)
|   |— flink_local_model_manager.py # Handles per-device
local training logic
|   |— flink_utils.py              # Helper functions for
state mgmt & aggregation
|       |— requirements.txt         # Flink-related
dependencies
|
|— models/
|   |— local/
|       |— model_template.py       # Feedforward NN
architecture for local nodes
|       |— device_001_model.pkl     # Example stored model
|       |— device_002_model.pkl
|       |— ...
|   |— global/
|       |— global_model.py          # Defines the global
model & aggregation logic
|       |— aggregator.py            # Aggregates weights
from local models
|       |— global_model_weights.pkl # Saved global model
weights

```

```

|   |   └─ global_update_scheduler.py           # Runs every X minutes
to update the global model
|   └─ utils/
|       └─ model_utils.py                       # Shared functions for
model loading/saving
|
└─ global_server/
    └─ app.py                                   # FastAPI or Flask
server for handling updates
|   └─ endpoints/
|       └─ receive_update.py                   # Endpoint: receive
local model weights
|       └─ get_global_model.py # Endpoint: return the latest global
model
|       └─ ...
|   └─ scheduler/
|       └─ aggregate_job.py # Periodic job that triggers model
aggregation
|   └─ config.json                             # Server configurations
|   └─ requirements.txt # Python dependencies for global server
|
└─ spark/
    └─ spark_batch_analysis.py # Batch analysis using full dataset
    └─ spark_streaming_analysis.py # Stream evaluation using Kafka
data
|   └─ spark_global_evaluator.py # Evaluates global model performance
|   └─ requirements.txt # Spark-related dependencies
|
└─ storage/
    └─ mongodb_init.py # Initialize MongoDB collections
    └─ mongodb_connection.py # Connection manager for MongoDB
    └─ schemas/
        └─ device_data_schema.json
        └─ local_model_schema.json
        └─ global_model_schema.json
        └─ predictions_schema.json
    └─ mongo_config.json # DB configuration (host, port, db name)
|

```

```
|— visualization/
|   |— superset_setup.sh                # Setup Apache Superset
|   |— dashboards/
|   |   |— per_device_dashboard.json
|   |   |— global_model_metrics.json
|   |   └─ comparisons_dashboard.json
|   └─ superset_config.py              # Connects Superset to MongoDB
|
|— orchestration/
|   |— run_all.sh                      # Starts the full pipeline end-to-end
|   |— start_streaming.py              # Starts Kafka producer
|   |— start_flink.py                  # Starts Flink job
|   |— start_global_server.py          # Starts global aggregation server
|   |— start_spark_jobs.py             # Starts Spark jobs (batch +
streaming)
|   └─ stop_all.sh                    # Stops all running services
|
|— config/
|   |— project_config.yaml             # General project-level settings
|   |— paths.json                     # Defines directories used in project
|   |— scheduler_config.json          # Defines timing for global aggregation
|   └─ environment.env                # Environment variables
|
|— utils/
|   |— logger.py                      # Centralized logging system
|   |— metrics.py                     # Functions for evaluating model accuracy
|   |— helpers.py                     # Generic utility functions
|   └─ time_utils.py                  # Functions for time formatting and scheduling
|
|— README.md                          # Full documentation and setup instructions
|— requirements.txt                    # Top-level dependencies
└─ LICENSE                            # Open-source license (MIT or Apache 2.0)
```

Technologies Used

- Python (scripts, ML models)
 - Apache Kafka (streaming)
 - Apache Flink (stream processing, local models)
 - Neural networks (feedforward, backpropagation)
 - Global model server (Python/Flask or FastAPI)
 - Apache Spark (batch & streaming analytics)
 - MongoDB (data storage)
 - Superset (visualization)
-

Conclusion

This project establishes a **full end-to-end IoT data streaming and ML analytics system**, combining **real-time streaming, federated local model training, global model aggregation, and heavy analytics**, all integrated with **MongoDB storage and BI visualization**. The architecture is **scalable, modular, and designed for continuous improvement** as new IoT devices or datasets are added.